

EEL4930/EEL5840 Fall 2016 – Project 3

Multilayer Perceptron (MLP) – MNIST Dataset

November 24, 2016

Due: December 10, 2016, 11:59 PM

This project will be graded with a letter grade with respect to presentation (25%), methods (35%) and results (40%).

In this project you should **train a Multilayer Perceptron (MLP) on the MNIST dataset**, which is a database of the **10 handwritten digits**. The MNIST dataset (28×28) contains **60k images for training** and **10k images for testing**. The **last 10k images in the training dataset** should be used as validation dataset for cross-validation, which include **early stopping and hyper parameter choices**.

NOTE: If your computer is slow, you can use less exemplars for training, but please **specify how large is the training set** because this will affect the quality of the results. The **size of the cross-validation** and test sets should remain at **10k**.

You should address (at least) the following topics:

1. **Download the MNIST dataset** and code from the course website (courtesy of <http://ufldl.stanford.edu/>).
2. Load images and labels using **loadMNISTImages** and **loadMNISTLabels** scripts. For example:

```
Training_images = loadMNISTImages('train-images-idx3-ubyte');  
Training_labels = loadMNISTLabels('train-labels-idx1-ubyte');
```
3. Train multilayer perceptrons (MLPs) for classification. We suggest using **one and two hidden layers only**, otherwise training will be too long. You have to **experiment with the number of units in each layer**. You can also **experiment with dropout** (i.e. with a certain probability set to zero weights during training, which may improve the generalization of the network).
4. We suggest to **downsample the image to decrease the number of inputs** or use **PCA to project to a subspace**. You should experiment with the **level of downsampling** and also with the **dimensionality of the subspace**.
5. Use samples **50000 to 60000** from the training dataset as the **validation dataset** for **early stopping**.

6. Report results on the test set as a *confusion matrix* (i.e. show the **true labels** and the **assigned classes** on a 10×10 matrix (one for each digit). If **the classifier is perfect you should have only the diagonal populated**, so **classification errors will appear in the off-diagonal entries**). Discuss the **results you obtain** and state if they **make sense or not**.
7. Show **learning curves** comparing **training and validation cost function**. Do the same for the **accuracy**.
8. **Compare performance** with Stochastic Gradient Descent (SGD) and SGD-momentum.

Hints:

- Use **ReLU activation function**, $f(net) = \begin{cases} 0 & net \leq 0 \\ net & net > 0 \end{cases}$, for **better results and faster convergence**.

The project requires a report explaining the experimental procedures you followed and you must include data to support your conclusions. Please use the format of an IEEE Transactions paper (limited to 7 double column pages). You can download the format from the IEEE website. This means you have to write a brief intro to the theory, explain well the methods and present carefully the results and conclusions. Remember that any scientific paper should, by definition, contain sufficient information such others can replicate your results. A scientific paper must also contain ORIGINAL material only. If you happen to use text or equations from other source you have to reference what you cut and paste (this is not allowed in a normal publication, but here it is OK provide you reference). Of course, we expect the results to be done by the student alone.